AI / Agents

Rational agent: perceives environment + uses knowledge, history, goals to choose action maximizing performance.
Limited rationality: exact optimal action may be too expensive.
AI view: acting rationally $\neq$ acting human-like.

Uninformed Search

Alg	C	O	Key / Cost
BFS	Y	unit	shallowest; $T, S = O(b^d)$
UCS	Y	Y	min g; $O(b^{C^*/\epsilon})$
DFS	finite m	N	deepest; $T = O(b^m), S = O(bm)$
DLS	$d \leq l$	N	cutoff l; $T = O(b^l), S = O(bl)$
IDS	Y	unit	repeated DLS; $T = O(b^d), S = O(bd)$

b: branch; d: optimal depth; m: max depth; ϵ: min step cost.
Trap: BFS optimal only for equal costs. UCS orders by total path cost. DFS may fail in infinite depth.

BFS / UCS / DFS Mini

BFS: FIFO fringe, append children to end. UCS: priority queue by accumulated $g(n)$. DFS: stack/LIFO, low memory.
Mini: $S \to A = 2, S \to B = 1, A \to G = 2, B \to G = 10$. UCS expands B first, but returns S-A-G cost 4.

Informed Search

Greedy best-first: minimizes $h(n)$ only. Fast if heuristic good; not optimal.
A*: minimizes total estimated path cost.

$$f(n) = g(n) + h(n)$$

g: cost so far. h: estimated remaining cost.
Admissible: $h(n) \leq h^*(n)$, never overestimates.
Consistent: $h(n) \leq c(n, n') + h(n')$.
Consistent $\Rightarrow$ admissible. Admissible $\not\Rightarrow$ consistent.
A* optimal if h admissible. Optimally efficient if h consistent.
Dominance: if $h_2 \geq h_1$ and both admissible, h_2 expands fewer/equal nodes.
Relaxed problem: remove constraints; relaxed solution cost gives admissible h.

A* Mini

Edges: $S \to A = 2$, $S \to B = 1$, $A \to G = 4$, $B \to G = 10$; $h(A) = 3, h(B) = 1, h(G) = 0$.
$f(A) = 2 + 3 = 5$, $f(B) = 1 + 1 = 2$. A* expands B first, sees G cost 11, then expands A and returns S-A-G cost 6.
Trap: In A*, goal generated $\neq$ goal selected. Return when goal is popped.

Local Search

Hill climbing: keep one current state; move to best neighbor. Not complete/optimal.
For minimization: accept child if $v(child) < v(current)$.
Problems: local optimum, plateau, ridge.
Fix: random restart; sideways moves with visit limit; macro moves; limited lookahead.
Beam search: keep best k states per level. Low memory; not complete/optimal.
Trap: Beam width k is not branching factor b. Discarded states are gone.

Simulated Annealing / GA

Simulated annealing: random neighbor; accept better always; accept worse with probability:

$$p = \exp\left(\frac{v(current) - v(next)}{T}\right)$$

for minimization. High T: exploration; low T: greedy.
Cooling too fast $\Rightarrow$ local optimum. Too slow $\Rightarrow$ inefficient.
Mini: $v(c) = 10, v(n) = 11, T = 2$. Worse move: $p = e^{-1/2} = 0.607$. If random $r = 0.4$, accept.
Genetic algorithm: population + fitness + selection + crossover + mutation.

$$P_i = \frac{f(S_i)}{\sum_j f(S_j)}$$

Selection gives uphill pressure; crossover/mutation gives exploration. Encoding matters.

Games

Minimax: 2-player, deterministic, perfect-info, zero-sum. MAX assumes MIN optimal.

$$V(s) = \begin{cases} U(s), & terminal \\ \max_{s'} V(s'), & MAX \\ \min_{s'} V(s'), & MIN \end{cases}$$

Time $O(b^m)$, space $O(bm)$.
Mini: MAX choices: A $\to \{3, 5\}$, B $\to \{2, 9\}$. MIN backs A=3, B=2, MAX chooses A.
Alpha-beta: same result as minimax; skips irrelevant branches.
α: best lower bound for MAX. β: best upper bound for MIN.
Prune when $\alpha \geq \beta$. Best $O(b^{m/2})$, worst $O(b^m)$. Ordering matters.
Mini: Root MAX has left value 100, so $\alpha = 100$. At right MIN, first child 20 gives $\beta = 20$. Since $\beta \leq \alpha$, prune rest.
Expectiminimax: chance node:

$$V = \sum_i p_i V_i$$

Mini: Chance children 2,4 with $p = .5$: $V = 3$.

ML Basics

Supervised: labelled data. Classif. = discrete label; regression = numeric value.
Unsupervised: no labels. Clustering = group similar examples.
Reinforcement: feedback/reward. **Association**: co-occurring patterns.
Train set: fit model. Valid. set: tune/prune/early stop. Test set: final comparison only.

Evaluation

$$Acc = \frac{TP + TN}{TP + TN + FP + FN}$$

$$Error = 1 - Acc$$

$$P = \frac{TP}{TP + FP}, \quad R = \frac{TP}{TP + FN}, \quad F_1 = \frac{2PR}{P + R}$$

Precision: predicted positive, how often right? Recall: real positive, how many found?
Use F_1 when classes imbalanced or FP/FN both matter.
Stratification: preserve class proportions. **k-fold CV**: each fold test once; average metrics.
LOOCV: $k = n$; expensive; deterministic. **Paired t-test**: same splits; fold-wise differences.
Trap: Never tune on test set.

Clustering Concepts

Unsupervised partition: same-cluster similar, different-cluster dissimilar.
Good clustering = high cohesion + high separation.
Centroid = mean point, may not be real item. Medoid = actual central item.

$$c_i = \frac{1}{|K_i|} \sum_{x \in K_i} x$$

Single link:

$$d(K_a, K_b) = \min_{x \in K_a, y \in K_b} d(x, y)$$

Complete link:

$$d(K_a, K_b) = \max_{x \in K_a, y \in K_b} d(x, y)$$

Average link: average pairwise distance.
Single link: chain effect. Complete link: compact clusters.
Davies-Bouldin:

$$DB = \frac{1}{k} \sum_i \max_{j \neq i} \frac{S_i + S_j}{d(c_i, c_j)}$$

Small DB = compact + far apart = good.
Mini: $K_1 = \{0, 2\}$, $K_2 = \{8, 10\}$. $c_1 = 1, c_2 = 9, S_1 = S_2 = 1$, DB = .25.

K-Means

Partitional, iterative, distance-based clustering.
Procedure: choose k centroids $\to$ assign to nearest $\to$ recompute centroids $\to$ repeat until stable.
Objective:

$$SSE = \sum_{i=1}^{k} \sum_{x \in K_i} d(x, c_i)^2$$

Time $O(tkmn)$; space $O((m + k)n)$.
Pros: simple, popular, efficient, usually fast.
Cons: need k; init-sensitive; local minima; outliers; poor on non-spherical clusters.
Mini: Points $1, 2, 8, 9$, $k = 2$, init $c_1 = 1, c_2 = 9$. Clusters $\{1, 2\}, \{8, 9\}$; new centroids 1.5, 8.5; stable.

Image Compression

1024^2 image, 2×2 blocks $\Rightarrow 512^2$ vectors, dimension 4:

$$rate = \frac{k \cdot 4 \cdot 8 + 512 \cdot 512 \log_2 k}{1024 \cdot 1024 \cdot 8}$$

NN Clustering / Hierarchical

Nearest Neighbor Clustering: incremental; process each item once.
If $d(new, K) \leq t$, join closest cluster; else new singleton. Usually single-link:

$$d(x, K) = \min_{y \in K} d(x, y)$$

Order-sensitive; threshold-sensitive; $O(n^2)$; chain effect.
Mini: $t = 2$: A starts; B dist 1 joins; C min 2 joins; D min 1 joins; E min 3 forms new cluster.
Hierarchical: dendrogram. No need k in advance.
Agglomerative = bottom-up merge. Divisive = top-down split.
Time $O(n^3)$ or $O(n^2 \log n)$; space $O(n^2)$.
Trap: Update distance matrix after each merge.

kNN

Lazy, instance-based. Training = store data. Query = compute distances to all examples, choose k nearest.
Classif. = majority vote. Regression = average.
Distances:

$$D_E = \sqrt{\sum_i (a_i - b_i)^2}$$

$$D_M = \sum_i |a_i - b_i|$$

Normalization:

$$a_i = \frac{v_i - \min v_i}{\max v_i - \min v_i}$$

Weighted kNN:

$$w_i = \frac{1}{D(x, x_i)}, \quad \hat{y} = \sum_i w_i y_i$$

Nominal: same 0, different 1. Missing: worst-case / ignore depending setting.
Small $k \Rightarrow$ overfit/noisy. Large $k \Rightarrow$ smooth/biased.
Curse of dim.: distances become similar; irrelevant attrs hurt.
Mini: Training $(4, L), (9, M), (8, M), (15, H), (7, M)$, query 5. Distances $1, 4, 3, 10, 2$. 1NN=L; 3NN=L,M,M $\Rightarrow$ M.

1R

One-rule classifier using one attribute.
For each attr: for each value predict majority class; count errors; choose attr with lowest train error.
Numeric: sort values; candidate breakpoints where adjacent class changes:

$$breakpoint = \frac{v_i + v_{i+1}}{2}$$

Mini: humidity: high = 3 yes,4 no $\Rightarrow$ no, errors 3; normal = 6 yes,1 no $\Rightarrow$ yes, errors 1; total 4.

Naive Bayes

Choose class with largest posterior score:

$$\hat{c} = \arg\max_c P(c) \prod_j P(x_j | c)$$

Assumption: attrs conditionally independent given class.
Bayes:

$$P(H|E) = \frac{P(E|H)P(H)}{P(E)}$$

Denominator common across classes, ignore for argmax.
Laplace:

$$P(x = v|c) = \frac{count(x = v, c) + 1}{count(c) + k}$$

m-estimate:

$$P(x = v_i|c) = \frac{n_i + mp_i}{n + m}$$

Gaussian numeric:

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/(2\sigma^2)}$$

Missing new attr: omit from product. Missing train value: omit from counts/denominator.
Mini: $P(Y) = 9/14$, $P(sunny|Y) = 2/9$, $P(cool|Y) = 3/9$. Score $Y = (9/14)(2/9)(3/9) = 1/21$.
Trap: Use class-conditional probabilities, not unconditional probabilities.

Decision Trees

Choose split with highest information gain.

$$H(S) = -\sum_i p_i \log_2 p_i$$

$$H(S|A) = \sum_v \frac{|S_v|}{|S|} H(S_v)$$

$$Gain(S, A) = H(S) - H(S|A)$$

Mini: 9 yes, 5 no: $H(S) = .940$. If $H(S|Outlook) = .693$, gain = .247.
Stop: pure; identical attrs but mixed labels $\to$ majority; empty subset $\to$ parent majority.
Overfit: train low, test high. Causes: small data, noise, too many branches.
Pre-prune: max depth / min leaf / min gain. Risk: underfit.
Post-prune: full tree + valid. set; replace subtree with leaf if valid. accuracy improves/same.
Rule pruning: tree $\to$ rules; remove harmless preconditions; sort by valid. accuracy.
Numeric attrs: sort values; candidate thresholds at class-change midpoints.
Missing values in DT gain: assign branch probabilities from non-missing examples.
Cost-sensitive split:

$$\frac{Gain(S, A)^2}{Cost(A)} \quad \text{or} \quad \frac{2^{Gain(S,A)} - 1}{(Cost(A) + 1)^w}$$

Gain ratio:

$$SplitInfo = -\sum_i \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}$$

$$GainRatio = \frac{Gain}{SplitInfo}$$

Gain ratio penalizes many-valued attrs.
Trap: Do not maximize remaining entropy. Maximize entropy reduction.

Perceptron

Single-layer neural classifier; linear boundary.

$$net = w \cdot p + b = \sum_i w_i p_i + b$$

$$a = step(net) = \begin{cases} 1, & net \geq 0 \\ 0, & net < 0 \end{cases}$$

$$e = t - a, \quad w_{new} = w_{old} + ep, \quad b_{new} = b_{old} + e$$

Decision boundary:

$$w \cdot p + b = 0$$

Converges iff data linearly separable. Cannot solve XOR / nonlinear boundary. No bias ⇒ boundary passes origin.
Mini: $w = [0, 0], b = 0, p = [0, 1], t = 0$. $a = 1$, $e = -1$, $w = [0, -1]$, $b = -1$.
Trap: Perceptron learning rule not guaranteed for non-linearly separable data.

MLP / Backprop

MLP with hidden layer + nonlinear differentiable activation forms nonlinear boundary. One hidden layer is universal approximator, but theorem does not tell how to find weights.
Sigmoid:

$$f(x) = \frac{1}{1 + e^{-x}}, \quad f'(x) = f(x)(1 - f(x))$$

Tanh derivative:

$$f'(x) = 1 - a^2$$

SSE:

$$E = \frac{1}{2} \sum_i (d_i - o_i)^2$$

Generic update:

$$\Delta w_{pq} = \eta \delta_q o_p, \quad w_{pq} \leftarrow w_{pq} + \Delta w_{pq}$$

Output neuron:

$$\delta_q = (d_q - o_q) f'(net_q)$$

Sigmoid output:

$$\delta_q = (d_q - o_q) o_q (1 - o_q)$$

Hidden neuron:

$$\delta_q = f'(net_q) \sum_i w_{qi} \delta_i$$

Sigmoid hidden:

$$\delta_q = o_q (1 - o_q) \sum_i w_{qi} \delta_i$$

Momentum:

$$\Delta w_{pq}(t) = \eta \delta_q o_p + \mu \Delta w_{pq}(t - 1)$$

η: learning rate; $\mu \in [0, 1)$: momentum.
Mini: $d = 1, o = .8, o_p = .5, \eta = .1$. $\delta = (1 - .8)(.8)(.2) = .032$. $\Delta w = .1(.032)(.5) = .0016$.

Backprop Procedure / Issues

Procedure: forward pass → compute output errors → backprop hidden errors → update weights → repeat.
Stopping: error threshold; max epochs; early stop on valid. set.
Input neurons: numeric directly; nominal via one-hot.
Output: binary → 1 output; k classes → k outputs. Often use .9/.1 instead of 1/0 to avoid flat gradient.
Hidden heuristic: roughly $(n_{in} + n_{out})/2$, then tune.
Too few hidden units ⇒ underfit. Too many ⇒ overfit/slow.
LR too small ⇒ slow. LR too large ⇒ oscillate/diverge.
Overfit causes: small data, noise, too many trainable params.
Fix: early stopping, different init, adaptive LR, dropout.
Local minima: GD may find local not global min.
Vanishing gradient: lower layers get tiny updates.
Trap: Backprop needs differentiable activation. Hidden neurons have no direct target values.

Deep / Autoencoder / CNN

Deep learning: multiple hidden layers; learns hierarchical features.
Autoencoder: target equals input:

$$y = x$$

Narrow hidden layer learns compressed representation.
Stacked autoencoder: unsupervised layer-wise pretraining, then supervised fine-tuning. Helps initialization/generalization.
CNN: for images/grid data; reduces parameters.
Local connectivity: local receptive fields.
Weight sharing/convolution: same filter across locations; feature map.

$$feature = \sum_{i,j} x_{ij} w_{ij}$$

Max pooling:

$$pool = \max(x_1, \ldots, x_r)$$

Pooling reduces dimension and gives shift invariance.
LCN:

$$x' = \frac{x - \mu}{\sigma}$$

Dropout: randomly deactivate neurons during training; reduces co-adaptation/overfit.

SVM

Maximum-margin hyperplane. SVs define boundary.

$$w \cdot x + b = 0$$

$$w \cdot x + b = \pm 1$$

$$margin = \frac{2}{\|w\|}$$

Hard margin:

$$\min \frac{1}{2} \|w\|^2$$

$$s.t. \quad y_i (w \cdot x_i + b) \geq 1$$

Dual / classify:

$$w = \sum_i \lambda_i y_i x_i$$

$$f(z) = \sum_i \lambda_i y_i (x_i \cdot z) + b$$

$$\hat{y} = \text{sign}(f(z))$$

Soft margin:

$$0 \leq \lambda_i \leq C$$

Large C: fewer train errors, narrower margin. Small C: wider margin, more violations.
Kernel:

$$K(u, v) = \Phi(u) \cdot \Phi(v)$$

$$K(x, y) = (x \cdot y + 1)^p$$

$$K(x, y) = e^{-\|x - y\|^2 / (2\sigma^2)}$$

Kernel trick: compute dot product in feature space without explicit mapping. Linear separator in feature space = nonlinear boundary in original space.
Mercer: valid kernel symmetric, continuous, positive semidefinite.
Mini: $w = (1, 1), z = (2, 1), b = 0$: $f = 3 > 0 \Rightarrow +1$.
Trap: Not all points define boundary; SVs do.

SVM / Perceptron / NN

Perceptron: simple, efficient, linear only.
Backprop NN: nonlinear, expressive, but many params, slow, local minima, architecture/LR sensitive.
SVM: margin + kernels, strong generalization, but kernel/C tuning and train cost matter.
Linear SVM: optimal linear boundary in original space.
Nonlinear SVM: optimal linear boundary in feature space.

Ensembles

Good ensemble = individual accuracy + error diversity.
Majority vote:

$$\hat{y} = mode(h_1(x), \ldots, h_M(x))$$

Weighted vote:

$$\hat{y} = \arg\max_c \sum_m \alpha_m \mathbf{1}[h_m(x) = c]$$

	Diversity	Train	Vote
Bag	bootstrap	parallel	equal
Boost	hard cases	seq.	weighted
RF	boot. + features	parallel	equal

Bagging reduces variance. Boosting focuses on difficult cases. RF lowers tree correlation.
Trap: Many identical classifiers do not help.

Bagging

Bootstrap aggregating. Sample n examples with replacement; train one classifier per sample; majority vote / average.
About 63% of original examples appear in a bootstrap sample.
Regression:

$$\hat{y} = \frac{1}{M} \sum_{m=1}^{M} h_m(x)$$

Useful for unstable learners e.g. DT. Easy to parallelize. Less useful for stable learners.
Mini: Original IDs 1,2,3,4,5. Bootstrap sample: 2,5,2,1,4.
Trap: Bagging samples with replacement and uses equal voting.

Boosting / AdaBoost

Sequential learners; later classifiers focus on examples earlier classifiers got wrong.
Weighted error:

$$\epsilon_t = \sum_i p_i e_i$$

$e_i = 1$ if misclassified.

$$\beta_t = \frac{\epsilon_t}{1 - \epsilon_t}$$

Correct examples:

$$p_i \leftarrow p_i \beta_t$$

then normalize. Classifier vote weight:

$$\alpha_t = \log \frac{1}{\beta_t}$$

Mini: 10 examples weight .1. Classifier gets 3 wrong: $\epsilon = .3$, $\beta = .3/.7 = .429$. Correct examples downweighted; wrong examples relatively upweighted.
Cons: sensitive to noisy labels/outliers; sequential not parallel.
Trap: Final boosting vote is weighted, not equal.

Random Forest

RF = bagging + random feature selection + unpruned DTs.
For each tree: bootstrap sample; grow DT without pruning; at each split choose k random features; select best among k using IG; majority vote.
Rule:

$$k \approx \log_2 m$$

m = total features.
Performance depends on high tree strength and low correlation.
Increasing k: strength ↑, correlation ↑.
Pros: often better than single DT; resists overfit; faster than AdaBoost with similar accuracy.
Trap: Standard RF trees are not pruned; do not use all features at every split.

Probability Basics

Conditional:

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

Product:

$$P(A, B) = P(A|B)P(B) = P(B|A)P(A)$$

Independence:

$$P(A, B) = P(A)P(B)$$

Disjunction:

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

Bayes:

$$P(H|E) = \frac{P(E|H)P(H)}{P(E)}$$

Total probability:

$$P(X) = \sum_i P(X|Y_i)P(Y_i)$$

Chain rule:

$$P(x_1, \ldots, x_n) = \prod_i P(x_i | x_1, \ldots, x_{i-1})$$

Normalize:

$$P(X|e) = \alpha P(X, e)$$

Mini: $P(M) = 1/50000, P(S|M) = .5, P(S) = 1/20$. $P(M|S) = .5(1/50000)/(1/20) = .0002$.
Trap: Do not reverse $P(A|B)$ and $P(B|A)$.

Bayesian Networks

BN = directed acyclic graph + CPTs.

$$P(x_1, \ldots, x_n) = \prod_i P(x_i | Parents(x_i))$$

Node cond. independent of non-descendants given parents.
Markov blanket = parents + children + children's parents.
Node independent of all outside blanket given blanket.
If Boolean vars have at most k parents: $O(n 2^k)$ params instead of $O(2^n)$.
Naive Bayes = special BN: class node parent of all features.
CPT learning by counting:

$$P(B) = \frac{\#b}{\#b + \#\neg b}$$

$$P(A|B, E) = \frac{\#a}{\#a + \#\neg a}$$

within rows where B, E fixed.
Mini: $P(j, m, a, \neg b, \neg e) = P(j|a)P(m|a)P(a|\neg b, \neg e)P(\neg b)P(\neg e)$.
Trap: BN must be acyclic. Missing edge means stated cond. independence, not arbitrary absolute independence.

Exact Inference

Goal: posterior such as $P(X|E = e)$.
Enumeration:

$$P(X|e) = \alpha P(X, e) = \alpha \sum_h P(X, e, h)$$

Variable elimination: same math, but multiply factors and sum out hidden vars one at a time, reusing intermediate factors.
Example:

$$P(B|j, m) = \alpha \sum_e \sum_a P(B)P(e)P(a|B, e)P(j|a)P(m|a)$$

Mini: Scores $b = .00059224, \neg b = .0014919$. Normalize $\Rightarrow P(B|j, m) = (.284, .716)$.
Trap: Do not eliminate query/evidence vars. Always normalize. Elimination order affects factor size.

Classifier Comparison

Alg	Exam summary
kNN	simple; no training; slow query; high memory; scale/irrelevant attrs hurt; append easy.
DT	interpretable; efficient; overfits; unstable; usually rebuild for updates.
NB	fast; one scan; robust; independence assumption; count update easy.
SVM	margin/kernel model; costly training; kernel/C tuning; re-train hard.
NN	nonlinear; expressive; slow; local minima; architecture/LR sensitive.
RF	accurate; robust; less interpretable; many trees raise prediction cost.

Common Exam Traps

A* returns goal when selected, not generated.
BFS optimal only for equal costs. UCS uses accumulated path cost.
Greedy uses h only; A* uses $g + h$.
Admissible ≠ consistent; consistent ⇒ admissible.
Precision denom. = predicted positives; recall denom. = real positives.
Valid. set for tuning/pruning/early stopping; test set for final evaluation.
Normalize numeric features for kNN/K-means.
Naive Bayes multiplies class-conditional probabilities.
DT information gain uses weighted child entropy.
Backprop needs differentiable activation.
Perceptron cannot solve XOR.
SVM boundary defined by SVs.
Bagging samples with replacement; boosting is sequential.
BN inference requires normalization.